Anomaly Detection in Traffic Scenarios: A Benchmarking Approach Using Pre-Trained AI Models

Author: Nicholas Jiang

Abstract:

The rapid advancement of autonomous vehicle technologies has underscored the importance of effective traffic accident detection systems. Current systems, however, often rely on labor-intensive hand-labeled datasets and make assumptions about fixed camera perspectives and static backgrounds. In this paper, we propose a novel approach to traffic accident detection that utilizes a pre-trained AI model as a benchmark for optimal driving behavior. By comparing real-world driving behaviors with those generated by the AI model, we can measure deviations and identify potential accidents when deviations exceed a certain threshold. Our methodology aims to address the limitations of existing systems by enabling the analysis of first-person video data from moving vehicles in an unsupervised manner.

Introduction:

As autonomous vehicles continue to become more prevalent in our transportation systems, ensuring safety on the roads is of paramount importance. Traffic accident detection is a crucial component of this safety framework, enabling rapid response to incidents and aiding in the prevention of secondary accidents. Traditional methods for traffic accident detection often involve the use of fixed cameras at specific locations, such as intersections or busy roads, or require labor-intensive hand-labeled datasets for training machine learning models. These approaches, while effective in certain scenarios, have notable limitations. Fixed cameras only cover a limited area and may miss incidents occurring outside their field of view. Hand-labeled datasets, on the other hand, are costly to create and maintain, and their utility is dependent on the quality and completeness of the annotations.

In this paper, we propose an innovative approach to traffic accident detection that addresses these limitations by leveraging the capabilities of modern autonomous driving AI models. These models, trained on extensive real-world data, exhibit expert-level driving behavior and can provide a useful benchmark for normal driving activity. We postulate that by comparing the behavior of human drivers to the "optimal" behavior of these AI models, we can detect anomalies indicative of potential accidents.

Specifically, our approach involves measuring the deviation of each vehicle's behavior from the AI model's behavior in real-world dashcam footage. If the deviation exceeds a certain threshold, we predict a potential accident. This method allows for the unsupervised analysis of first-person video data from moving vehicles, making it adaptable to a wide range of driving scenarios.

While still in its conceptual stage, we believe our approach presents a promising avenue for future research and development in the field of traffic accident detection. In the following sections, we will delve into the details of our proposed methodology, discuss its potential implications, and address possible challenges in its implementation.

Related Work:

A significant piece of related work in the field of traffic accident detection is the study titled "Unsupervised Traffic Accident Detection in First-Person Videos" by Junaid Ahmed Ghauri et al. This paper presents an unsupervised approach to accident detection based on predicting the future locations of traffic participants.

The authors' model, which they implemented using the PyTorch framework, leverages the power of variational autoencoders (VAEs) and generative adversarial networks (GANs). In their approach, they utilize VAE-GANs to predict the future locations of traffic participants based on past and current observations. The system identifies anomalies when the predictions are inconsistent or when the accuracy of the predictions suddenly drops, indicating a potential accident.

A key strength of their methodology is its unsupervised nature, which eliminates the need for labor-intensive hand-labeled datasets. Additionally, their system operates effectively with first-person video data from moving vehicles, overcoming the limitations of static camera positions.

The authors evaluated their model using a new dataset (A3D) they introduced, along with the publicly available HEV-I dataset. Their model significantly outperformed the baselines, demonstrating the effectiveness of their unsupervised anomaly detection approach in identifying traffic accidents.

While this work presents a robust approach to accident detection, it is primarily focused on predicting the future locations of traffic participants and detecting anomalies based on these predictions. In contrast, our proposed method extends the concept of anomaly detection by leveraging a pre-trained AI model as a benchmark for optimal driving behavior. We aim to detect anomalies by comparing the real-world driving behaviors with those generated by the AI model, thereby introducing a novel perspective in the field of traffic accident detection.

Methodology:

Our approach for anomaly detection in traffic scenarios combines the use of a pre-trained AI model, feature extraction from dashcam footage, and a measure of deviation from normal driving behavior. Here, we describe our proposed solution in detail.

1. **Pre-Trained AI Model**: The first step in our approach involves the selection of a pre-trained AI model that exhibits expert-level driving behavior. This model should be capable of understanding complex traffic scenarios and making safe and efficient driving decisions. Ideally, this model would have been trained on a large, diverse dataset of real-world driving scenarios, enabling it to generalize well to unseen situations. For our purposes, the model serves as a benchmark for "optimal" driving behavior.
2. **Data Collection**: We gather dashcam footage from a variety of driving scenarios. This footage serves as our primary dataset for anomaly detection. It should ideally include a diverse range of scenarios, including normal driving, aggressive driving, and accident situations. The use of dashcam footage enables our system to analyze driving scenarios from a first-person perspective, similar to the perspective of an autonomous vehicle.
3. **Feature Extraction**: For each vehicle in the dashcam footage, we use computer vision techniques to extract a set of features that represent the current driving scenario. These features could include the positions and velocities of nearby vehicles, road conditions, traffic signals, and any other relevant information. The extracted features should correspond to the input requirements of the pre-trained AI model.
4. **Benchmarking Using the AI Model**: With the extracted features, we simulate the decisions that the pre-trained AI model would make in the same scenario. These decisions represent the "optimal" driving behavior according to the AI model.
5. **Calculating Deviation**: We calculate the deviation between the actual behavior of the vehicle in the footage and the "optimal" behavior suggested by the AI model. This deviation serves as a measure of how far the vehicle's behavior deviates from the expected norm.
6. **Crash Prediction**: We predict a potential crash when the calculated deviation exceeds a certain threshold. This threshold can be determined empirically based on the performance of the system on a validation dataset.

Our methodology offers a unique approach to traffic accident detection by comparing real-world driving behaviors against a benchmark of optimal driving behavior. By identifying significant deviations from this benchmark, we aim to detect anomalies that could indicate potential accidents. This approach allows for the analysis of first-person video data from moving vehicles in an unsupervised manner, overcoming some of the limitations of existing methods.

Conclusion and Future Work:

In this paper, we have proposed a novel approach to traffic accident detection that leverages a pre-trained AI model as a benchmark for optimal driving behavior. By measuring the deviation of real-world driving behavior from this benchmark, we aim to detect anomalies that could indicate potential accidents. Our methodology is designed to address some of the limitations of existing systems, offering a way to analyze first-person video data from moving vehicles in an unsupervised manner.

While our approach is still in the conceptual stage, we believe it holds significant promise for improving traffic accident detection. The use of a pre-trained AI model as a benchmark provides a unique perspective on what constitutes "normal" driving behavior, and the ability to measure deviations from this benchmark could offer a powerful tool for identifying potential accidents.

Looking forward, we see several areas for further exploration and development. Firstly, the validation of our approach will be crucial. This will involve implementing a prototype of our system and testing its performance on a diverse dataset of real-world driving scenarios. Through such testing, we can refine the method used to calculate deviation, adjust the threshold for anomaly detection, and identify any areas where the performance of the system could be improved.

Secondly, we plan to explore ways to address the issue of perspective difference. As we noted earlier, most AI models for autonomous driving are trained on first-person perspective data, whereas our system is designed to analyze the behavior of other vehicles based on footage captured from our vehicle's perspective. Techniques such as data augmentation and multi-view training could potentially be used to mitigate this issue.

Lastly, we envision that our approach could be expanded to include other types of anomalous behavior beyond just accidents. For example, the system could potentially be adapted to detect aggressive driving, traffic violations, or other forms of unsafe behavior. Such expansions could further enhance the utility of our system and contribute to overall road safety.

In conclusion, we believe that our proposed approach represents a promising new direction in the field of traffic accident detection. We look forward to further developing this idea and exploring its potential to contribute to safer, more efficient transportation systems.

References:

Yao, Y., Xu, M., Wang, Y., Crandall, D. J., & Atkins, E. M. (2019). Unsupervised traffic accident detection in first-person videos. *2019 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. https://doi.org/10.1109/iros40897.2019.8967556